MogScan: Real-Time, Multi-Person, Explainable Facial-Aesthetics Scoring on a Commodity Webcam

Final Project — Startup / Product Track (Modified YC Application)

Ben Blair
Cyril Sokolov
AI Course, Fall 2026

Abstract

Face-rating ("looksmaxxing") apps are some of the most downloaded AI consumer products of the last two years. Omoggle, a webcam "mog battle" site, recently spread across Twitch. Current products, however, rate one face at a time, give a number without a credible explanation, usually upload the user's selfie to a server, and do not show that their scores agree with real human judgements. We built **MogScan**, a browser app that watches a USB webcam and gives *every* person in the frame (up to six) a live PSL score (1–8), a community tier (Sub5 … HTN, Chadlite, Chad), and a ranked list of the measured facial proportions that raise or lower that score: canthal tilt, philtrum-to-chin ratio, FWHR, and twelve more. Processing runs fully on-device. The pipeline frontalizes MediaPipe's 478-landmark face mesh and computes 15 metrics. It then combines an explainable, sign-constrained ridge model with a fine-tuned MobileNetV3. On the official SCUT-FBP5500 60/40 split, our fused model reaches Pearson $r = 0.917$ against the mean of 60 human raters. That is higher than the dataset's published ResNeXt-50 baseline ($r = 0.8777$), and the model runs in real time in a browser tab. We also report subgroup bias, which "looksmax rules" do and do not match human ratings, and a user study with **[run study]** participants.

1 Introduction

Motivation. "Looksmaxxing" is the online subculture of measuring and improving one's appearance. It has turned into a large consumer AI market. Industry reporting put the leading face-rating app, Umax, at over 7 million downloads and roughly $500K per month in subscription revenue by mid-2024 [19]. In 2026 *Omoggle*, an Omegle-style site that pairs two strangers on webcam and declares a "mogger", spread across Twitch after streamers such as xQc and Asmongold played it on stream [6]. The demand is clear. People want to know where they stand on the community's PSL scale and, more importantly, *why*.

The products serving this demand have four recurring problems: (1) **one face at a time**: selfie apps rate one photo, and Omoggle is strictly 1v1; (2) **unexplained numbers**: the scores are opaque, or the "feature analysis" is generic text; (3) **privacy**: most apps upload a face photo to a server; (4) **no validation**: we found no competitor that reports how well its scores agree with human raters. Coverage of Omoggle notes that camera angle and distance easily flip its results [6].

What we built. MogScan is a web app. Plug in any USB webcam, pick it from the camera menu, and everyone in frame is scored live (Figure 1). For each person MogScan shows:

[Insert figures/hero.png: MogScan Snapshot with 2–3 consenting people in frame, Guides + Mog battle on]

Figure 1: MogScan scoring several people at once from a single USB webcam. Each face gets a tracked ID, a PSL score and tier, and overlay guides for the measurements behind it (canthal-tilt lines, facial thirds, jaw and cheekbone widths). The side panel lists each person's top strengths and weaknesses with the measured value, the average of the top-20%-rated faces, and the looksmax community's reference range.

- a **PSL score** (1–8), the percentile, and the community tier (Sub3, Sub5, LTN, MTN, HTN, Chadlite, Chad, Adamlite+; female variants LTB/MTB/HTB, Stacylite, Stacy);
- the **top three strengths and weaknesses**, each with its measured value, its effect on the score in PSL points, the average of top-rated faces, and the community's reference range;
- all **15 measurements** on demand, plus on-video guides showing where they were taken.

A *Mog battle* mode ranks everyone in frame. A Kind/Brutal toggle changes the wording (Strength vs. "Halo", Room to grow vs. "Failo"). Snapshot saves a shareable image, and Record exports an anonymized numeric session log for our user study. Nothing leaves the device.

Contributions.

(1) A real-time, multi-person, fully client-side pipeline: landmarks, tracking, pose gating, frontalization, metrics, two models, and fusion, running at camera frame rate in a browser tab.
(2) An *explainable* geometric model whose per-feature attributions have signs that agree with the evidence, obtained with a sign-constrained ridge regression. We measure the accuracy cost of this constraint.
(3) An evaluation on SCUT-FBP5500: accuracy against human raters, subgroup bias, and a data-driven check of popular looksmax rules.
(4) Evidence from real users: a usability and value study (§5.4).

2 Modified YC Application

Describe what your company does in 50 characters or less. Live, explainable face ratings for groups.

What is your company going to make? A webcam product that rates every face in view in real time and explains the score with measured facial proportions. The first product is the free web app described in this paper. Next come a *streamer overlay* (an OBS browser source for Twitch and Kick creators who already run mog-battle content) and a mobile app with progress tracking.

Why did you pick this idea? Why now? Two things came together in 2026. First, demand: Omoggle showed that group and live face-rating is entertainment people will watch and stream [6]. Second, capability: browser ML (MediaPipe Tasks [14], ONNX Runtime Web) is now fast enough to run a 478-point face mesh for six people *and* a CNN on a laptop, without a server. That makes "on-device, nothing uploaded" free for us to offer and costly for server-based competitors to match.

Who are your users and how do you know they need it? (a) 18–25-year-olds in the looksmaxxing community, who already pay $4–$10 per week for selfie raters [12, 20]; (b) friend groups and parties, where multi-person scoring is a social game; (c) streamers. Evidence from our own study is in §5.4: SUS **[run study]**, NPS **[run study]**, **[run study]**% found the group mode fun, and **[run study]**% said they would pay for an extended version.

Who are your competitors? What do you understand that they don't? See Table 1. Our insight is that *explanations are the product.* Users do not just want a number. They want to know *which* features help or hurt them. Explanations only earn trust if they agree with the evidence, so we validate the score against human raters and build the explanation model so that its signs cannot contradict the data (§4.5).

How will you make money? Free live web app as the acquisition channel. **MogScan Pro** at $4.99/month adds history, progress tracking, a side-profile scan, and a "glow-up plan" of evidence-based, non-surgical suggestions. That is cheaper than the $4–$10 *per week* the incumbents charge [12]. **Creator license** at $9.99/month: the OBS overlay with branding and battle leaderboards. Our marginal compute cost per user is close to zero because inference runs on the user's device.

Traction so far. Working product (this paper), an accuracy result above the dataset's published baseline, and a user study with **[run study]** participants (§5.4).

What is the biggest risk? Ethical and reputational harm: body-image damage, use on minors, and rating people without consent (§6). A practical risk is data licensing: SCUT-FBP5500 is licensed for non-commercial research, so a commercial release needs a model retrained on licensed or self-collected ratings.

3 Related Work

3.1 Competing products

Omoggle [18] is the closest product and the inspiration for this project. It runs a 468-landmark face mesh in the browser and turns 10-second 1v1 duels into an ELO ladder (Molecule … Chadlite, Chad, Slayer). Its scoring is a black box, and press coverage notes that camera angle, distance and props easily change results [6]. MogScan addresses this with *pose gating* and *frontalization* (§4.3). The selfie apps (Umax, LooksMax AI, and dozens of PSL raters) show strong demand, but they put results behind paywalls and upload photos [12, 19, 20]. A parental-monitoring review rates PSL apps 18+ and warns about obsessive checking [1]; we take that seriously in §6. Qoves Studio sells human-written facial-aesthetics reports at a much higher price point [22].

3.2 Research

Facial beauty prediction (FBP). Early work regressed attractiveness on hand-crafted geometric ratios and found that a few proportions explain a large share of the variance in ratings [7, 9, 13]. SCUT-FBP [23] and SCUT-FBP5500 [16] made deep FBP standard. On the 60/40 split, the SCUT-FBP5500 paper reports Pearson correlations of 0.8298 (AlexNet), 0.8513 (ResNet-18) and 0.8777 (ResNeXt-50), and 0.8997 for ResNeXt-50 under 5-fold cross-validation. Recent large hybrid models report around 0.92 [2]. These works target offline accuracy on single photos. We target real-time, multi-person, *explained* predictions on a webcam, with a model small enough for a browser.

Psychology of attractiveness. Meta-analyses find strong agreement between raters on attractiveness within and across cultures [15]. That agreement is what makes a learned predictor possible. Averageness, symmetry and sexual dimorphism are the best-supported cues [21]. Facial width-to-height ratio (FWHR) is linked to perceived dominance [4]. Looksmax heuristics such as "positive canthal tilt" or a "1:2 philtrum-to-chin ratio" mostly come from forums, not studies. We test them against human ratings directly (Figure 4).

Face tracking. We use MediaPipe's FaceLandmarker (478 landmarks including irises, plus a facial transformation matrix) [8, 14] and MobileNetV3 [11] for the CNN.

Explainability and fairness. Additive attributions such as SHAP [17] decompose a prediction into per-feature contributions. Our geometric model is additive by construction, so its attributions are exact. Commercial face analysis has documented subgroup disparities [3], and debiasing for aesthetic prediction has been proposed [5]. We report per-subgroup error in Table 4.

4 Implementation

4.1 Design process and ideation

Our initial plan was a direct Omoggle-style design: a face mesh, looksmax ratios, and a score from hand-picked "ideal" ranges for each ratio. Four findings during development reshaped it:

(1) **Hand-picked ideals are arbitrary.** Their weights had no justification, and some community ideals do not match the measurements MediaPipe's landmarks produce. We switched to *learning* the weights from 5,500 faces rated by 60 people [16].
(2) **Head pose leaks into every metric.** On the dataset, our symmetry metric correlated $r = 0.97$ with |yaw|, so it was measuring head turn, not the face. Frontalization reduced this to $r = 0.67$, and pose gating in the app handles the rest (§4.3).
(3) **Unconstrained explanations contradict the data.** The metrics are correlated with each other (e.g. jaw : cheekbone width

Table 1: Competitors and alternatives. Prices come from public listings and third-party reports (Oct. 2026) and vary by region and listing.

	Omoggle [6, 18]	**UMAX / LooksMax AI** [12, 19]	**PSL rating apps** [1, 20]	**Qoves Studio** [22]	**MogScan (ours)**
Input	Live webcam, 1v1 with a stranger	One selfie	One selfie	Uploaded photos	Live webcam, **up to 6 people**
Output	0–10 score, win/loss, ELO rank	Overall + sub-scores (jaw, skin, …)	PSL score + tier	Written report	PSL, tier, percentile, 15 measured metrics with PSL effect
Explanation	Mentions symmetry, canthal tilt, jaw; none shown per user	Category scores	Category scores	Human-written analysis	Per-feature attribution from a validated, sign-constrained model
Validated vs. human raters	Not reported; reportedly sensitive to angle and props	Not reported	Not reported	Not reported	Pearson $r = 0.917$ on a held-out test split
Privacy	Claims on-device	Server upload	Server upload	Server upload	On-device, nothing stored
Price	Free	≈$4–10 / week	Subscription	≈$20–200 one-time	Free; Pro $4.99 / month (planned)

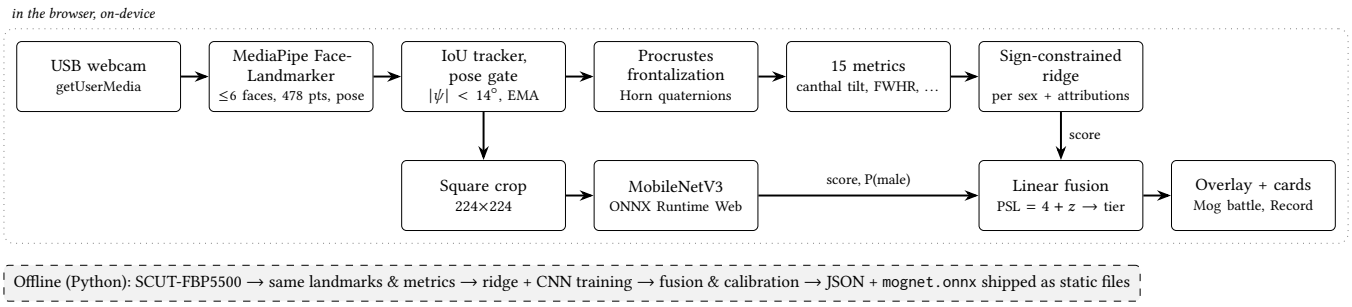


Figure 2: System architecture. The browser pipeline (top) and the offline training pipeline (bottom) share the same landmark model, crop rule and metric code. The TypeScript metric implementation is checked against the Python reference in unit tests.

vs. jaw angle, $r = -0.75$; FWHR vs. face length : width, $r = -0.68$). An ordinary ridge regression therefore gave FWHR a large *positive* weight (+0.26 for male and +0.46 for female faces, in rating units per SD), even though FWHR on its own correlates *negatively* with ratings ($r \approx -0.25$). Users with a high FWHR would have been told it was a strength while the raw data says the opposite. We now constrain the signs (§4.5).

(4) **Camera pose ≠ apparent pose.** Our first pose gate used MediaPipe's transformation matrix. It rejected a frontal face near the edge of a multi-person frame, because the matrix includes the off-axis viewing angle. We now gate on the apparent pose from frontalization, which is what the metrics actually depend on.

For the interface we chose a browser app over a native OpenCV window. A browser tab runs on any laptop with any USB camera, can be shared as a link, and keeps everything on-device. Live cards update five times per second, while the overlay is drawn on every frame.

4.2 Landmarks, tracking and pose gating

MediaPipe FaceLandmarker runs in `VIDEO` mode on the GPU delegate, falling back to CPU, with up to six faces, and returns 478 normalized 3D landmarks per face. We read the *apparent* head pose, yaw ψ and pitch, from the rotation that aligns the face to our frontal template (§4.3). We do not use MediaPipe's facial transformation matrix for this. That matrix measures rotation relative to the camera ray, so in testing a perfectly frontal face near the edge of a wide frame read as $-17°$ of yaw. The Procrustes yaw correlates $r = 0.99$ with MediaPipe's yaw on the centered dataset images, at about 1.2× the scale. A greedy IoU tracker (match threshold 0.25, 1 s time-to-live) keeps a stable ID per person. Metrics are smoothed with an exponential moving average ($\alpha = 0.15$) and updated *only* while the face is near-frontal ($|\psi| < 14°$, $|\text{pitch}| < 20°$). When someone turns away, the card says "face the camera" and the score is frozen. This is a direct fix for the angle sensitivity reported for Omoggle.

4.3 Frontalization

Landmarks are converted to pixel-space 3D (xW, yH, zW), centered, and scaled to unit RMS radius. We then find the rotation that best aligns them to a mean frontal template using Horn's closed-form

Table 2: Metrics (MediaPipe landmark indices in parentheses).

Metric	Definition
Canthal tilt	mean angle of inner→outer canthus vs. pupil line (133–33, 362–263); + = outer higher
Eye spacing (ESR)	IPD (468–473) / bizygomatic width (234–454)
FWHR	bizygomatic width / height glabella–upper lip (9–0)
Midface ratio	pupil line–upper lip height / IPD
Lower third	(subnasale–menton) / (glabella–menton) (2, 152, 9)
Chin : philtrum	(lower lip–menton) / (subnasale–upper lip) (17–152, 2–0)
Nose : intercanthal	alar width (129–358) / intercanthal width (133–362)
Mouth : nose	mouth width (61–291) / alar width
Jaw : cheekbone	bigonial (172–397) / bizygomatic width
Lower : upper lip	lower-lip height (14–17) / upper-lip height (0–13)
Eye openness	mean lid aperture / eye width (159–145, 386–374)
Brow–eye distance	brow (105, 334) to upper lid, / IPD
Face length : width	forehead–menton (10–152) / bizygomatic width
Asymmetry	mean mirror error of 13 bilateral pairs about the midline, / IPD
Jaw angle	projected angle zygion–gonion–menton (234–172–152, 454–397–152)

quaternion method [10]: the top eigenvector of a 4×4 symmetric matrix built from the cross-covariance. The template is a generalized-Procrustes mean of 1,257 near-frontal *training* faces. Metrics are computed on the x, y coordinates of the aligned shape, after removing residual roll so that the line between the pupils is horizontal. The browser uses a Jacobi eigen-solver port of the same algorithm. Unit tests check it against the Python reference within 10^{-3} on five faces, including the most strongly turned faces in the dataset.

4.4 The 15 metrics

All lengths are normalized by the interpupillary distance (IPD) or the bizygomatic width, so the metrics do not depend on how far someone is from the camera. Table 2 gives the landmark definitions.

4.5 Explainable geometric model

For each sex we standardize every metric (z_i, clipped to ±4) and fit

$$\hat{y}_{\text{geo}} = b + \sum_{i=1}^{15} \left(\beta_i z_i + \gamma_i (z_i^2 - 1)\right) + \text{pose terms},$$

where the pose terms ($|\psi|$, |pitch − median|) absorb any remaining pose effects during training and are set to zero at inference. In the *sign-constrained* variant, $\text{sign}(\beta_i)$ must match the sign of metric i's marginal correlation with the rating, and $\gamma_i \le 0$, so any curvature can only express an "ideal range" (an inverted U). We solve it as non-negative ridge regression on a re-signed design matrix, with α chosen by 5-fold cross-validation. Because the model is additive, the contribution $c_i = \beta_i z_i + \gamma_i (z_i^2 - 1)$ is exact and has zero mean over the training set, so it reads as "points above or below an average face." The UI shows c_i/σ_y in PSL units and the top three positive and negative contributions. The "top-tier average" shown for each metric is that metric's mean over the top 20% of rated training faces of the same sex.

4.6 CNN

We fine-tune an ImageNet-pretrained MobileNetV3-Large [11] with a two-output head (beauty score, sex logit; loss = smooth-L1 + 0.2×BCE). Crops are square boxes 1.35× the landmark bounding box, black-padded, resized to 224×224. The browser uses exactly the same rule, so training and live inputs are framed the same way. Training used 85% of the official training split, with 15% held out for validation. Settings: AdamW, one-cycle learning rate with a peak of 3×10^{-4}, batch 32, 25 epochs, flips, ±8° affine and color jitter. We kept the epoch with the best validation correlation ($r = 0.915$). The model is exported to ONNX (opset 17, 13.9 MB) and runs in ONNX Runtime Web (WASM). A round-robin scheduler sends one face to the CNN every 120 ms, and the outputs are smoothed per track.

Table 3: SCUT-FBP5500 official 60/40 split (test $n = 2200$). Target: mean rating of 60 raters (1–5).

Model	Pearson r ↑	MAE ↓	RMSE ↓
AlexNet [16]	0.8298	–	–
ResNet-18 [16]	0.8513	–	–
ResNeXt-50 [16]	0.8777	–	–
Ours: geometry only (constrained ridge)	0.674	0.392	0.501
Ours: MobileNetV3 CNN	0.914	0.210	0.278
Ours: fused (deployed)	**0.917**	**0.208**	**0.273**

4.7 Fusion and PSL calibration

A linear regression on the validation set combines the two models: $\hat{y} = 1.00\,\hat{y}_{\text{cnn}} + 0.09\,\hat{y}_{\text{geo}} + b$. We map to PSL with $\text{PSL} = 4 + (\hat{y} - \mu)/\sigma$, where $\mu = 2.99$ and $\sigma = 0.69$ are the mean and SD of the human ratings, clipped to $[1, 8]$. With this mapping the community tier cut-offs land on familiar percentiles: MTN covers the middle (3.75–4.5), HTN is roughly the top 31%, Chadlite the top 10% (≥ 5.25), and Chad the top 2.3% (≥ 6). The detected sex chooses the geometric model and the tier names, and users can correct it by clicking the sex badge on their card.

4.8 Interface and privacy

The app is built with React and TypeScript on Vite. Everything is drawn on one canvas: the video, corner-bracket boxes colored by tier, the PSL and tier label, and optional guides and mesh. All model files and wasm runtimes are served from the app itself, so it also works offline on demo day. No frame is ever sent over the network; Snapshot and Record save files locally. Record logs *numbers only* (per-person PSL, model outputs, yaw, fps). An opening screen asks for confirmation that everyone in frame is 18+ and has consented.

4.9 Replicating the system

All code is in our repository. `README.md` lists the nine commands that rebuild every model and number in this paper from the public dataset, and `npm run dev` starts the app. Every number in this document is generated from the result files by `make_figures.py`.

5 Evaluation

5.1 Agreement with human raters

Table 3 and Figure 3 show the results. The fused model reaches $r = 0.917$ on the held-out test split. That is higher than all three of the dataset paper's 60/40 baselines, although it is below recent large research models (≈0.92 [2]), which are not designed for real-time use in a browser. The geometry-only model ($r = 0.674$) is much weaker than the CNN. Proportions explain a large part of

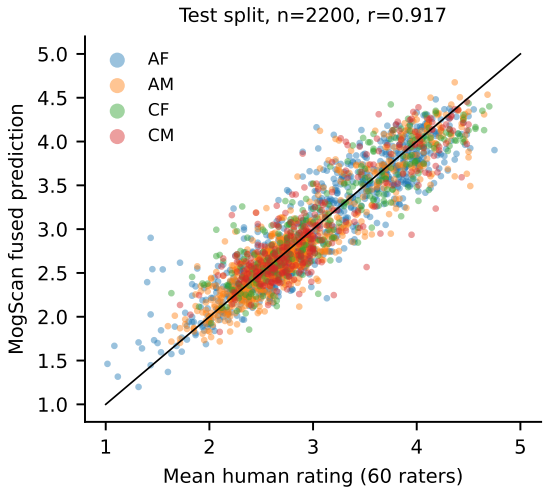


Figure 3: Fused prediction vs. mean human rating on the test split, colored by subgroup (A/C = Asian/Caucasian, F/M).

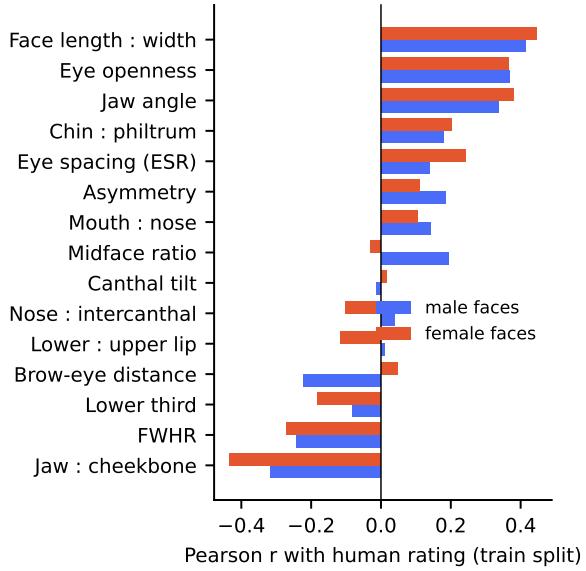


Figure 4: Correlation of each frontalized metric with human ratings (training split), split by sex.

attractiveness, but skin, hair, expression and other texture cues matter too. Its value is in the explanations. The sex head reaches 99.1% accuracy. On the PSL scale the mean absolute error is 0.30 points, the exact tier matches in 61.7% of test faces, and the tier is within one step in 98.1%.

5.2 Do looksmax rules hold up?

Figure 4 shows each metric's raw correlation with the human rating, measured on frontalized landmarks. The strongest single metric is *Face length : width* ($r = 0.43$, averaged across sexes). Several community rules do *not* hold up for these raters:

- **Canthal tilt**, the best-known looksmax metric, has essentially no correlation ($r \approx 0.00$). Before frontalization it seemed to matter ($r \approx 0.28$), but that came from head pose.
- **"Hunter eyes":** more *open* eyes go with higher ratings ($r \approx 0.37$).
- **Wide jaw and high FWHR:** faces that are *narrower* in the jaw relative to the cheekbones ($r \approx -0.38$), with lower FWHR ($r \approx -0.26$) and softer projected jaw angles ($r \approx 0.36$), are rated higher. That is a V-shaped face, the opposite of the "wide, square jaw" ideal.
- **Symmetry** looks *negatively* related to attractiveness on the full dataset, but the effect disappears ($r \approx 0.00$) for faces within 3° of frontal. It comes from residual pose, since head-turned photos received higher ratings ($r = 0.28$ for |yaw|, male faces).

A longer chin relative to the philtrum ($r \approx 0.19$) and wider-set eyes ($r \approx 0.19$) agree with community advice. These results describe one dataset and one rater panel, and East Asian beauty norms may favour V-shaped faces. Still, they show why MogScan displays the *data-derived* target next to the community range instead of repeating forum rules, and why pose control is necessary for any measurement-based rating.

Table 4: Fused model by subgroup (test split). Signed bias = mean(pred − human).

Group	*n*	*r*	MAE	Signed bias
Asian female	792	0.918	0.216	-0.005
Asian male	811	0.928	0.184	-0.032
Caucasian female	285	0.905	0.236	-0.007
Caucasian male	312	0.881	0.225	-0.008

5.3 Subgroup bias

Table 4 breaks the error down by the dataset's four subgroups. SCUT-FBP5500 has only Asian and Caucasian faces, rated by one panel of 60 young raters. **People of other ethnicities are not represented at all**, so the model's behaviour on them is unvalidated. That is a serious limitation for a consumer product. The app's opening screen says so, and collecting a more diverse dataset is our first item of future work.

5.4 User study

[The user study has not been run yet. Follow study/protocol.md with 8–15 participants, export the Google Form to study/responses.csv and the session logs to study/logs/, run analyze_survey.py, analyze_logs.py and make_figures.py, then rebuild this PDF. Every number in this section is filled in automatically. Do not type numbers by hand.]

5.5 Performance and stability

[Fill from study logs (median PSL std. while still, median fps).] In our own development testing (Chrome on an Apple-silicon MacBook, photo mode, one face) the full pipeline ran at about 27–36 frames per second.

6 Ethics, Risks and Limitations

Rating faces can hurt people. Parental-safety reviewers rate PSL apps 18+ and warn that they can reinforce body dissatisfaction [1]. Our mitigations:

- an 18+ and consent screen;
- on-device processing with no storage;
- "Kind" wording by default;
- honest framing in the app ("a model's estimate of how one panel judged faces, not your worth");
- a wellbeing item in our survey.

Technical limitations:

- SCUT faces are studio-like frontal photos, while webcam frames vary in lighting and lens distortion, so there is domain shift;
- the mesh starts at the upper forehead, so true facial thirds and full face length are approximations;
- the jaw angle is a 2D projection;
- the dataset covers two ethnic groups and only binary sex labels.

Commercial use would require licensed training data.

7 Future Work

- **Diverse, licensed data:** collect consented ratings across many ethnicities, ages and rater groups, retrain, and publish subgroup error before any commercial launch.
- **Side-profile scan:** gonial angle, chin projection and the Ricketts E-line need a profile view. Prompt users to turn their head and fuse the measurements.
- **Streamer overlay:** an OBS browser source with a transparent background and battle ELO, aimed at the creators who popularized Omoggle.
- **Progress tracking and glow-up plan:** longitudinal, pose-normalized measurements with evidence-based, non-surgical suggestions (grooming, skin care, posture). Never surgery prompts.
- **Better CNN explanations:** Grad-CAM heatmaps for the texture cues the geometric model cannot see, and a WebGPU backend for higher frame rates with six people in frame.

References

[1] Bark. 2026. PSL App Review. https://www.bark.us/app-reviews/apps/psl-app-review/. Accessed Oct. 2026.

[2] Djamel Eddine Boukhari. 2025. SynergyNet: Fusing Generative Priors and State-Space Models for Facial Beauty Prediction. *arXiv preprint arXiv:2509.17172* (2025).

[3] Joy Buolamwini and Timnit Gebru. 2018. Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. In *Proceedings of the Conference on Fairness, Accountability and Transparency (PMLR, Vol. 81)*. 77–91.

[4] Justin M. Carré and Cheryl M. McCormick. 2008. In Your Face: Facial Metrics Predict Aggressive Behaviour in the Laboratory and in Varsity and Professional Hockey Players. *Proceedings of the Royal Society B* 275, 1651 (2008), 2651–2656.

[5] Michael Danner, Thomas Weber, Leping Peng, Tobias Gerlach, Xueping Su, and Matthias Rätsch. 2021. Ethically Aligned Deep Learning: Unbiased Facial Aesthetic Prediction. *arXiv preprint arXiv:2111.05149* (2021).

[6] Dexerto. 2026. What is Omoggle? The AI Face-rating Platform Taking over Twitch. https://www.dexerto.com/twitch/what-is-omoggle-the-ai-face-rating-platform-taking-over-twitch-3360363/.

[7] Yael Eisenthal, Gideon Dror, and Eytan Ruppin. 2006. Facial Attractiveness: Beauty and the Machine. *Neural Computation* 18, 1 (2006), 119–142.

[8] Ivan Grishchenko, Artsiom Ablavatski, Yury Kartynnik, Karthik Raveendran, and Matthias Grundmann. 2020. Attention Mesh: High-fidelity Face Mesh Prediction in Real-time. *arXiv preprint arXiv:2006.10962* (2020).

[9] Hatice Gunes and Massimo Piccardi. 2006. Assessing Facial Beauty through Proportion Analysis by Image Processing and Supervised Learning. *International Journal of Human-Computer Studies* 64, 12 (2006), 1184–1199.

[10] Berthold K. P. Horn. 1987. Closed-form Solution of Absolute Orientation Using Unit Quaternions. *Journal of the Optical Society of America A* 4, 4 (1987), 629–642.

[11] Andrew Howard, Mark Sandler, Grace Chu, Liang-Chieh Chen, Bo Chen, Mingxing Tan, Weijun Wang, Yukun Zhu, Ruoming Pang, Vijay Vasudevan, Quoc V. Le, and Hartwig Adam. 2019. Searching for MobileNetV3. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*.

[12] HyScaler. 2026. Looksmax AI: The Complete 2026 Guide. https://hyscaler.com/insights/looksmax-ai-self-improvement-app-men/. Accessed Oct. 2026.

[13] Amit Kagian, Gideon Dror, Tommer Leyvand, Isaac Meilijson, Daniel Cohen-Or, and Eytan Ruppin. 2008. A Machine Learning Predictor of Facial Attractiveness Revealing Human-like Psychophysical Biases. *Vision Research* 48, 2 (2008), 235–243.

[14] Yury Kartynnik, Artsiom Ablavatski, Ivan Grishchenko, and Matthias Grundmann. 2019. Real-time Facial Surface Geometry from Monocular Video on Mobile GPUs. *arXiv preprint arXiv:1907.06724* (2019).

[15] Judith H. Langlois, Lisa Kalakanis, Adam J. Rubenstein, Andrea Larson, Monica Hallam, and Monica Smoot. 2000. Maxims or Myths of Beauty? A Meta-analytic and Theoretical Review. *Psychological Bulletin* 126, 3 (2000), 390–423.

[16] Lingyu Liang, Luojun Lin, Lianwen Jin, Duorui Xie, and Mengru Li. 2018. SCUT-FBP5500: A Diverse Benchmark Dataset for Multi-Paradigm Facial Beauty Prediction. In *Proceedings of the International Conference on Pattern Recognition (ICPR)*. arXiv:1801.06345.

[17] Scott M. Lundberg and Su-In Lee. 2017. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*.

[18] Omoggle. 2026. Omoggle — 1v1 Mog Battle. https://www.omoggle.app/. Accessed Oct. 2026.

[19] Plisio. 2026. What Is Looksmax AI? Looksmaxxing, Glow-Up, and the Real Risks. https://plisio.net/ai/looksmax-ai. Cites Fortune (July 2024) for Umax download and revenue figures. Accessed Oct. 2026.

[20] Prodigy AI. 2026. Looksmaxxing – PSL Face Rating (App Store listing). https://apps.apple.com/app/id6477295133. Accessed Oct. 2026.

[21] Gillian Rhodes. 2006. The Evolutionary Psychology of Facial Beauty. *Annual Review of Psychology* 57 (2006), 199–226.

[22] SwitchTools. 2026. Qoves vs Dreamwave AI. https://www.switchtools.io/blog/qoves-vs-dreamwave-ai. Third-party price estimate. Accessed Oct. 2026.

[23] Duorui Xie, Lingyu Liang, Lianwen Jin, Jie Xu, and Mengru Li. 2015. SCUT-FBP: A Benchmark Dataset for Facial Beauty Perception. In *IEEE International Conference on Systems, Man, and Cybernetics (SMC)*.